

Mean episode reward

Total env steps (M)

4,096 agents $\times$ 128 steps/iter
 $\approx$ 524K transitions/update

PPO iteration

— Mean reward (EMA-20)
★ Selected checkpoint (iter 975)

